

```
In [4]: # setup
from IPython.core.display import display,HTML
display(HTML('<style>.prompt{width: 0px; min-width: 0px; visibility: collapse}</style>'))
display(HTML(open('rise.css').read()))

# imports
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
sns.set(style="whitegrid", font_scale=1.5, rc={'figure.figsize':(12, 8)})
```

CMPS 2200

Introduction to Algorithms

Overview

Today's agenda:

- Introductions
- Motivation for course
- Formalisms used throughout the course
- Navigating the course

What is an algorithm?

an explicit, precise, unambiguous, mechanically-executable sequence of elementary instructions, usually intended to accomplish a specific purpose.

-- Jeff Erickson

Examples?

BOB(n):

- for n down to 1
 - Sing “ i bottles of beer on the wall, i bottles of beer,”
 - Sing “Take one down, pass it around, $i - 1$ bottles of beer on the wall.”
- Sing “No bottles of beer on the wall, no bottles of beer,”
- Sing “Go to the store, buy some more, n bottles of beer on the wall.”

Lagniappe: Another old algorithmic drinking song is [The Barley Mow](#). I'll leave it to you to come up with the program to produce it.

Examples of algorithm-like things that are not algorithms?

BeAMillionaireAndNeverPayTaxes():

- Get a million dollars.
- If the tax man comes to your door and says, “You have never paid taxes!”
 - Say “I forgot.”

What makes a good algorithm?

- correct
- user-friendly
- many features
- robust
- simple
- secure
- low programmer cost
- **efficient**
 - runs quickly
 - requires little memory

Then, why study efficiency?

- separates feasible from infeasible
- correlates with user-friendliness

What if it took Google took 2 minutes to return results?

Simple warmup: **What does this do?**

```
In [26]: def my_function(a, b):
        for i,v in enumerate(a):
            if v == b:
                return i
        return -1
```

```
In [1]: def linear_search(mylist, key):
        """
        Args:
        mylist...a list
        key.....a search key
        Returns:
        index of key in mylist; -1 if not present
        """
        for i,v in enumerate(mylist):
            if v == key:
                return i
        return -1

linear_search([5,1,10,7,12,4,2], 12)
```

```
Out[1]: 4
```

What factors affect the running time of this algorithm?

- Input **size**
- Input **values**: is key at start or end?
- Hardware!
 - TI-85 vs. Supercomputer

We need a way to compare the efficiency of algorithms that abstracts away details of hardware and input.

Analysis of Linear Search, the long way

- Assign a time cost c_i to each line i .
- Figure out how often each line is run n_i
- total cost is the cost of each line multiplied by the number of times it is run

$\text{Cost}(\text{linear-search}, \text{mylist}, \text{key}) = \sum_i c_i * n_i$

```
In [28]: def linear_search(mylist, key):          # cost      number of times run
        for i,v in enumerate(mylist):        # c1         ?
            if v == key:                      # c2         ?
                return i                      # c3         ?
        return -1                           # c4         ?
```

Best/Average/Worst case

To deal with the effects of the input values on performance, we can consider three types of analysis:

- **Worst-case**: maximum time for any input of size n

```
linear_search([5,1,10,7,12,4,2], 9999)
```

- **Best case**: minimum time of any input of size n

```
linear_search([5,1,10,7,12,4,2], 5)
```

- **Average case**: expected time over all inputs of size n
 - Need some probability distribution over inputs

```
for (mylist, key) in ???:
    linear_search(mylist, key)
```

Worst-case analysis of linear search

Assume $n \leftarrow \text{len}(\text{mylist})$

```
In [29]: def linear_search(mylist, key):      # cost      number of times run
        for i,v in enumerate(mylist):        # c1         ?
            if v == key:                      # c2         ?
                return i                      # c3         ?
        return -1                           # c4         ?
```

$\text{Cost}(\text{linear-search}, n) = c_1n + c_2n + c_4$

Cost is now just a function of:

- input size n
- constants c (depend on machine, compiler, etc)

How granular should we get?

Consider this slightly different implementation:

```
In [30]: def new_linear_search(mylist, key):  # cost      number of times run
        for i in range(len(mylist)):          # c5         n
            if mylist[i] == key:              # c6         n
                return i                      # c3         0
        return -1                           # c4         1
```

$\text{Cost}(\text{new-linear-search}, n) = c_5n + c_6n + c_4$ (1)

$\text{Cost}(\text{linear-search}, n) = c_1n + c_2n + c_4$ (2)

There is a line between analyzing algorithms and optimizing a specific implementation (e.g., through profiling):



Big Idea: Asymptotic Analysis

- Ignore machine-dependent constants
- Focus on **growth** of running time
 - What happens in the limit as $n \rightarrow \infty$

$c_1n + c_2n + c_4 \approx c_5n + c_6n + c_4$

e.g., consider two algorithms with running times:

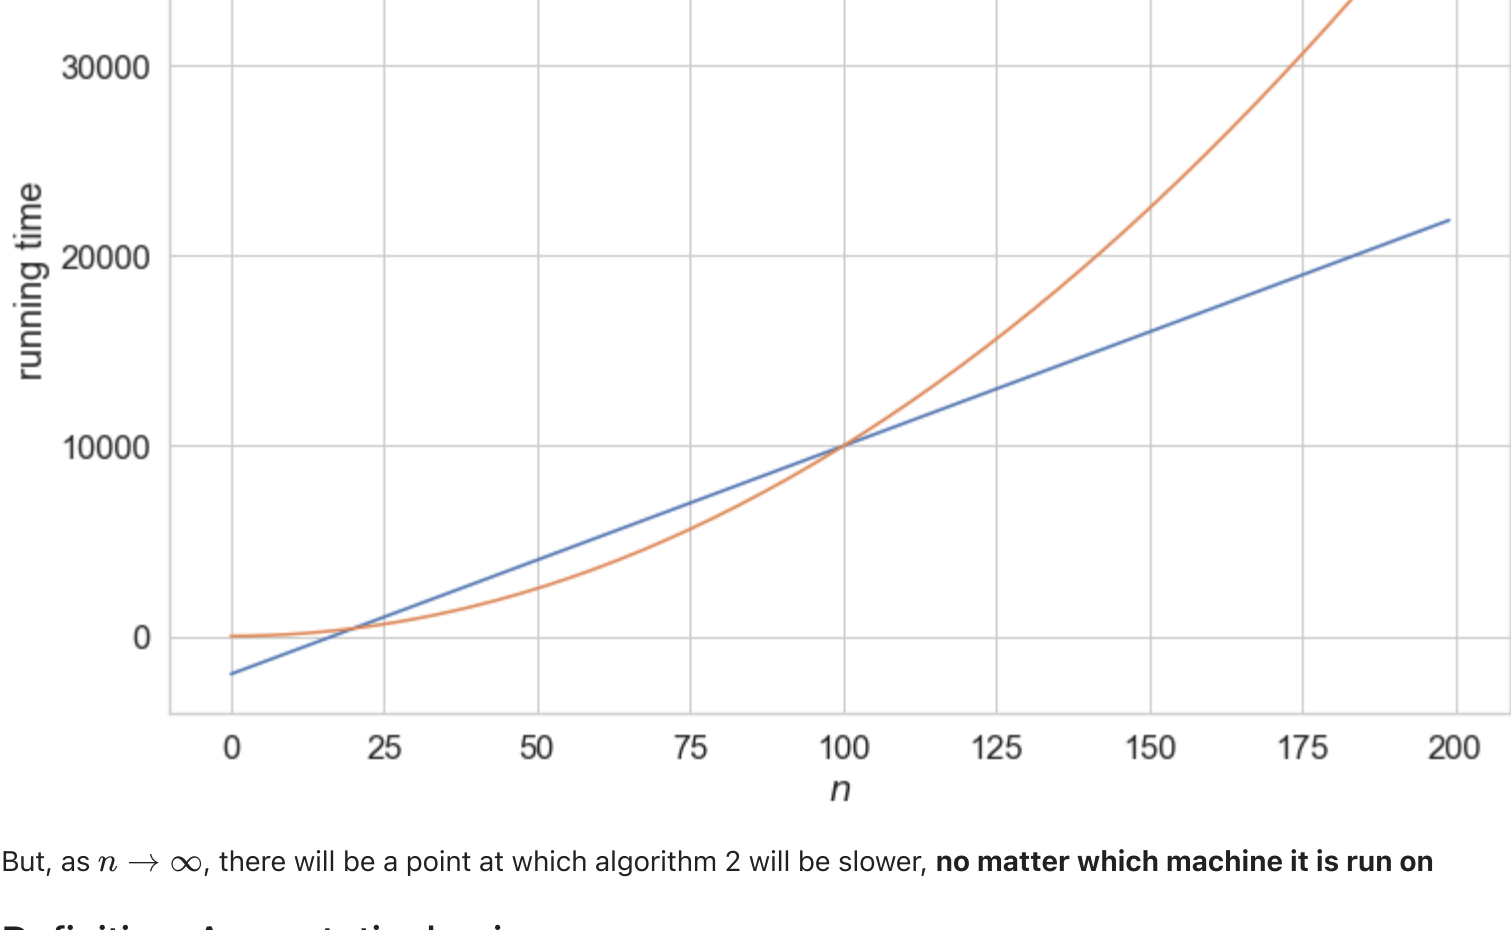
- algorithm 1: $c_1n + c_2$
- algorithm 2: $c_3n^2 + c_4n + c_5$

Depending on the machine-dependent constants, algorithm 2 may sometimes be faster than algorithm 1:

- algorithm 1: $120n - 2000$
- algorithm 2: n^2

```
In [5]: n = np.arange(200)
        timel = 120*n + - 2000
        time2 = n*n

        # plot
        plt.figure()
        plt.plot(n, timel, label='$120 n - 2000$')
        plt.plot(n, time2, label='$n^2$')
        plt.xlabel("$n$")
        plt.ylabel('running time')
        plt.legend()
        plt.show()
```



But, as $n \rightarrow \infty$, there will be a point at which algorithm 2 will be slower, **no matter which machine it is run on**

Definition: Asymptotic dominance

Function $f(n)$ **asymptotically dominates** function $g(n)$ if **there exist** constants c and n_0 such that

$g(n) \leq c \cdot f(n)$ for all $n \geq n_0$

e.g., n^2 asymptotically dominates $120n - 2000$

Proof:

Find c and n_0 such that

$120n - 2000 \leq c * n^2$ for all $n > n_0$

Let $c = 1$. Find an n_0 such that

$120n - 2000 \leq n^2$

for all $n \geq n_0$

$120n - 2000 \leq n^2$ (3)

$0 \leq n^2 - 120n + 2000$ (4)

$0 \leq (n - 100)(n - 20)$ (5)

When $n = 100$, $120n - 2000 = n^2$

For all $n \geq 100$, $120n - 2000 \leq n^2$

So, $c = 1$, $n_0 = 100$ satisfies the definition of asymptotic dominance.

```
In [32]: # show n_0
        n = np.arange(150)
        timel = 120*n + - 2000
        time2 = n*n

        # plot
        plt.figure()
        plt.plot(n, timel, label='$120 n - 2000$')
        plt.plot(n, time2, label='$n^2$')
        plt.axvline(100, color='g')
        plt.text(94,1000,'$n_0$', fontsize=18, color='g')
        plt.xlabel("$n$")
        plt.ylabel('running time')
        plt.legend()
        plt.show()
```



Asymptotic Notation

$\mathcal{O}(f(n)) = \{g(n) \mid f(n) \text{ asymptotically dominates } g(n)\}$ (6)

$\Omega(f(n)) = \{g(n) \mid g(n) \text{ asymptotically dominates } f(n)\}$ (7)

$\Theta(f(n)) = \mathcal{O}(f(n)) \cap \Omega(f(n))$ (8)

e.g.

$120n - 2000 \in \mathcal{O}(n^2)$

$10n^3 + 2n^2 - 100 \in \Omega(n^2)$

$14n^2 - 5n + 50 \in \Theta(n^2)$

We often abuse notation such as

$120n - 2000 = \mathcal{O}(n^2)$

or

$120n - 2000 \text{ is } \mathcal{O}(n^2)$

Analogy:

\mathcal{O}	Ω	Θ	o	ω
\leq	\geq	$=$	$<$	$>$

Course Overview

- Analyzing algorithms: methods to compute tight bounds on running time
- Designing algorithms: various approaches to designing efficient algorithms
 - lists, sequences, trees, graphs,...
- Distinct from typical courses like this, we will emphasize **parallel** algorithms from the start (next lecture)

Navigating the course

- Canvas: syllabus, dates, grades
- Diderot: interactive textbook
- Github: assignments, slides